

```
import numpy as np
import torch
import matplotlib.pyplot as plt
import torchvision
import torchvision.transforms as transforms
from torch import nn
from torchvision import utils
```

✓ Load the datasets

```
# Load MNIST dataset
```

```
train_data = torchvision.datasets.MNIST(root=".", train=True, transform=transforms
test_data = torchvision.datasets.MNIST(root=".", train=False, transform=transforms
```

```
➔ Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Failed to download (trying next):
<urlopen error [SSL: CERTIFICATE_VERIFY_FAILED] certificate verify failed: ce
```

```
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-images-idx3-ubyte.gz
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-images-idx3-ubyte.gz
100%|██████████| 9912422/9912422 [00:00<00:00, 15961007.53it/s]
Extracting ./MNIST/raw/train-images-idx3-ubyte.gz to ./MNIST/raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Failed to download (trying next):
<urlopen error [SSL: CERTIFICATE_VERIFY_FAILED] certificate verify failed: ce
```

```
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-labels-idx1-ubyte.gz
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-labels-idx1-ubyte.gz
100%|██████████| 28881/28881 [00:00<00:00, 491422.69it/s]
Extracting ./MNIST/raw/train-labels-idx1-ubyte.gz to ./MNIST/raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Failed to download (trying next):
<urlopen error [SSL: CERTIFICATE_VERIFY_FAILED] certificate verify failed: ce
```

```
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte.gz
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte.gz
100%|██████████| 1648877/1648877 [00:00<00:00, 4414755.70it/s]
Extracting ./MNIST/raw/t10k-images-idx3-ubyte.gz to ./MNIST/raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Failed to download (trying next):
<urlopen error [SSL: CERTIFICATE_VERIFY_FAILED] certificate verify failed: ce
```

```
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz
100%|██████████| 4542/4542 [00:00<00:00, 6477568.44it/s]
Extracting ./MNIST/raw/t10k-labels-idx1-ubyte.gz to ./MNIST/raw
```

```
# Hyperparameters
learning_rate = 0.001
epochs = 10
batch_size = 200
no_of_eigen_val = 30 # number of principal components (for PCA)

train_loader = torch.utils.data.DataLoader(dataset=train_data, batch_size=batch_si
test_loader = torch.utils.data.DataLoader(dataset=test_data, batch_size=len(test_d

# Creating a DataLoader for a subset of test_data
test_sample_loader = torch.utils.data.DataLoader(dataset=test_data.data[9705:9715

train_dataset = train_data.data.reshape(train_data.data.shape[0], train_data.data.
test_dataset = test_data.data.reshape(test_data.data.shape[0], test_data.data.shap
test_dataset_sampled = test_dataset[np.arange(9705, 9715), :] # slicing & creating
```

✓ Comparing PCA and autoencoders

```
# i/p data is a 2D tensor of shape (num_datapts, 784)
def PCA(input_data, top_k_ev):
    # mean calculation and centering
    input_mean = torch.mean(input_data, 0)
    centered_data_matrix = input_data - input_mean

    # Covariance Matrix
    cov_matrix = torch.matmul(centered_data_matrix.T, centered_data_matrix) # dot p

    # Eigenvalues and Eigenvectors
    eigen_values, eigen_vectors = torch.linalg.eigh(cov_matrix) # eigenvalues: vec
    sorted_eigen_vals, sorted_indices = torch.sort(eigen_values, descending=True) #

    # Select the top k eigenvectors and eigenvalues
    top_k_eigen_values, top_k_indices = sorted_eigen_vals[:top_k_ev], sorted_indice
    top_k_eigen_vectors = eigen_vectors[:, top_k_indices]

    assert top_k_eigen_vectors.shape == (784, top_k_ev) # checks the shape
    assert centered_data_matrix.shape == input_data.shape
    assert input_mean.shape == torch.Size([784])

    return top_k_eigen_vectors
```

```
pc = PCA(train_dataset.float(), no_of_eigen_val) # (784, top_k)

# reconstruct the original dataset (in its original high-dimensional space) using
def reconstruct_data(principal_components, dataset): # principal_components: matr

    # Create the projection matrix
    projection_matrix = torch.matmul(principal_components, principal_components.T)

    # project the dataset using the projection matrix
    projected_data = torch.matmul(dataset, projection_matrix)
    # each data point in dataset is projected back into the original 784-dimensiona

    assert projected_data.shape == (dataset.shape[0], principal_components.shape[0])
    return projected_data

reconstructed_test_data_sampled = reconstruct_data(pc, test_dataset_sampled.float)

# Autoencoder neural network
class AutoEncoder(nn.Module):
    def __init__(self):
        super(AutoEncoder, self).__init__()
        self.encoder = nn.Sequential( # encoder part
            nn.Linear(784,512), # i/p has 784 dimensions - 28x28
            nn.ReLU(),
            nn.Linear(512,256),
            nn.ReLU(),
            nn.Linear(256,128),
            nn.ReLU(),
            nn.Linear(128,30),
            nn.ReLU())
        self.decoder = nn.Sequential( # decoder part
            nn.Linear(30,128),
            nn.ReLU(),
            nn.Linear(128,256),
            nn.ReLU(),
            nn.Linear(256,784),
            nn.ReLU())

    # forward pass
    def forward(self,x): # shape of x (batch_size, 784))
        x = self.encoder(x)
        encoded_output = x
        x = self.decoder(x)
        return x, encoded_output

model1 = AutoEncoder()
criterion1 = nn.MSELoss() # loss function
```

```

optimizer1 = torch.optim.Adam(model1.parameters(), lr=learning_rate) # adam optim

# Training loop
training_loss=[]
for epoch in range(epochs):
    for images, labels in train_loader:
        images = images.reshape(images.shape[0], -1) # 2D to 1D vector
        outputs, _ = model1(images) # reconstructed images
        loss = criterion1(outputs, images) # reconstruction error
        training_loss.append(loss.item())

    optimizer1.zero_grad() # clears the gradients from the previous batch
    loss.backward() # backpropagation
    optimizer1.step() # update the model parameter


print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.item()))
#print("Epoch [{}/{}]: completed".format(epoch + 1, epochs))

Epoch [1/10]: completed, Loss: 0.0329
Epoch [2/10]: completed, Loss: 0.0259
Epoch [3/10]: completed, Loss: 0.0221
Epoch [4/10]: completed, Loss: 0.0205
Epoch [5/10]: completed, Loss: 0.0188
Epoch [6/10]: completed, Loss: 0.0173
Epoch [7/10]: completed, Loss: 0.0166
Epoch [8/10]: completed, Loss: 0.0172
Epoch [9/10]: completed, Loss: 0.0154
Epoch [10/10]: completed, Loss: 0.0160

plt.figure(figsize=(10, 6))
plt.plot(range(1,len(training_loss)+1),training_loss, color='green')
plt.title("Training Loss vs Iterations")
plt.xlabel("No: of Iterations")
plt.ylabel("Training Loss")
plt.grid()

```





```
# Model evaluation
model1.eval() # evaluation mode
with torch.no_grad(): # disables gradient calculation

    # reshaping and passing test images through the AE
    for images in test_sample_loader:
        images = images.reshape(10, 28*28) # reshape from 2D to 1D
        outputs, _ = model1(images.float()) # reconstructed images

# visualization and plotting
plt.rcParams["figure.figsize"] = (10,6) # rcParams:set global parameters for plot
for i in range (5): # iterates over 10 images
    fig, (ax1, ax2, ax3) = plt.subplots(1,3)

    ax1.imshow(images[i].detach().numpy().reshape(28,28), cmap = 'gray') # original
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs[i].detach().numpy().reshape(28,28), cmap = 'gray') # reconctr
    ax2.set_title('AE Reconstructed Image')
    ax2.axis("off")


    ax3.imshow(reconstructed_test_data_sampled[i].reshape(28,28), cmap = 'gray') # P
    ax3.set_title('PCA Reconstructed Image')
    ax3.axis("off")
    print("Reconstruction Error in AE:",np.dot(((images[i].detach().numpy())/255)-(o
                                                ((images[i].detach().numpy())/255)-(ou

    print("Reconstruction Error in PCA:",np.dot(((images[i].detach().numpy())/255)-(
                                                ((images[i].detach().numpy())/255)-(


plt.show()

Reconstruction Error in AE: 5.114889141803459
Reconstruction Error in PCA: 4.906910216143052
```


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 22.499340382207624

Reconstruction Error in PCA: 16.862376030476067


Original Image



AE Reconstructed Image



PCA Reconstructed Image



Reconstruction Error in AE: 15.27587834574187

Reconstruction Error in PCA: 16.0457551172948


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 13.645043006066087

Reconstruction Error in PCA: 10.714796514097081


Original Image



AE Reconstructed Image



PCA Reconstructed Image







Reconstruction Error in AE: 14.066630936640893

Reconstruction Error in PCA: 14.848193155147694


Original Image



AE Reconstructed Image



PCA Reconstructed Image



✓ Standard Autoencoder

```
# Under-complete AutoEncoder - one hidden layer
class Standard_AE(nn.Module):
    def __init__(self, hid):
        super(Standard_AE, self).__init__()
        self.hid = hid # stores the no: of neurons in the hidden layer (hid)
        self.encoder = nn.Sequential(
            nn.Linear(784, self.hid),
            nn.ReLU()
        )
        self.decoder = nn.Sequential(
            nn.Linear(self.hid, 784),
            nn.ReLU()
        )

    def forward(self, x):
```

```
x = self.encoder(x)
encoded_output = x
x = self.decoder(x)
return x, encoded_output
```

```
# Hiddensize 64
```

```
model_hid_64 = Standard_AE(64)
criterion_hid64 = nn.MSELoss()
optimizer_hid64 = torch.optim.Adam(model_hid_64.parameters(), lr=learning_rate)
```

```
training_loss_hid64 = []
```

```
# Training loop for the specified number of epochs
```

```
for epoch in range(epochs):
```

```
    for images, labels in train_loader:
```

```
        images = images.reshape(images.shape[0], -1)
```

```
        outputs, _ = model_hid_64(images) # reconstructed images
```

```
        loss = criterion_hid64(outputs, images)
```

```
        training_loss_hid64.append(loss.item())
```

```
    optimizer_hid64.zero_grad()
```

```
    loss.backward()
```

```
    optimizer_hid64.step()
```

```
print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.item()))
```

```
Epoch [1/10]: completed, Loss: 0.0182
```

```
Epoch [2/10]: completed, Loss: 0.0146
```

```
Epoch [3/10]: completed, Loss: 0.0134
```

```
Epoch [4/10]: completed, Loss: 0.0133
```

```
Epoch [5/10]: completed, Loss: 0.0125
```

```
Epoch [6/10]: completed, Loss: 0.0126
```

```
Epoch [7/10]: completed, Loss: 0.0125
```

```
Epoch [8/10]: completed, Loss: 0.0123
```

```
Epoch [9/10]: completed, Loss: 0.0112
```

```
Epoch [10/10]: completed, Loss: 0.0124
```

```
# Hiddensize 128
```

```
model_hid_128 = Standard_AE(128)
```

```
criterion_hid128=nn.MSELoss()
```

```
optimizer_hid128 = torch.optim.Adam(model_hid_128.parameters(), lr=learning_rate)
```

```
training_loss_hid128 = []
```

```
for epoch in range(epochs):
```

```
    for images, labels in train_loader:
```

```
        images=images.reshape(images.shape[0], -1)
```

```
        outputs, _ = model_hid_128(images)
```

```
        loss = criterion_hid128(outputs, images)
```

```
        training_loss_hid128.append(loss.item())
```



```
optimizer_hid128.zero_grad()
loss.backward()
optimizer_hid128.step()
```

```
print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.item()))
```

```
Epoch [1/10]: completed, Loss: 0.0122
Epoch [2/10]: completed, Loss: 0.0094
Epoch [3/10]: completed, Loss: 0.0085
Epoch [4/10]: completed, Loss: 0.0085
Epoch [5/10]: completed, Loss: 0.0078
Epoch [6/10]: completed, Loss: 0.0075
Epoch [7/10]: completed, Loss: 0.0081
Epoch [8/10]: completed, Loss: 0.0073
Epoch [9/10]: completed, Loss: 0.0074
Epoch [10/10]: completed, Loss: 0.0079
```

```
# Hiddensize 256
model_hid_256 = Standard_AE(256)
criterion_hid256=nn.MSELoss()
optimizer_hid256 = torch.optim.Adam(model_hid_256.parameters(),lr=learning_rate)
```

```
training_loss_hid256=[]
```

```
for epoch in range(epochs):
    for images,labels in train_loader:
        images=images.reshape(images.shape[0],-1)
        outputs, _ = model_hid_256(images)
        loss = criterion_hid256(outputs,images)
        training_loss_hid256.append(loss.item())
```

```
optimizer_hid256.zero_grad()
loss.backward()
optimizer_hid256.step()
```

```
print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.item()))
```


```
Epoch [1/10]: completed, Loss: 0.0097
Epoch [2/10]: completed, Loss: 0.0071
Epoch [3/10]: completed, Loss: 0.0062
Epoch [4/10]: completed, Loss: 0.0060
Epoch [5/10]: completed, Loss: 0.0057
Epoch [6/10]: completed, Loss: 0.0056
Epoch [7/10]: completed, Loss: 0.0050
Epoch [8/10]: completed, Loss: 0.0053
Epoch [9/10]: completed, Loss: 0.0053
Epoch [10/10]: completed, Loss: 0.0053
```

```
plt.plot(range(1,len(training_loss_hid64)+1),training_loss_hid64,label="Hidden Si
plt.plot(range(1,len(training_loss_hid64)+1),training_loss_hid128,label="Hidden S
```

```

plt.plot(range(1, len(training_loss_hid64)+1), training_loss_hid256, label="Hidden S
plt.legend()
plt.grid()
plt.title("Training Loss Vs Iterations")
plt.xlabel("No: of Iterations")
plt.ylabel("Training Loss")
plt.show()

```



```

model_hid_64.eval()
with torch.no_grad(): # disable gradient calculation
    for images in test_sample_loader:
        # print(images.shape)
        images = images.reshape(10, 28*28)
        outputs_hid64, _ = model_hid_64(images.float()) # Forward pass through the mo

```

```

model_hid_128.eval()
with torch.no_grad():
    for images in test_sample_loader:

```

```

for images in test_sample_loader:
    # print(images.shape)
    images = images.reshape(10,28*28)
    outputs_hid128, _ = model_hid_128(images.float())

model_hid_256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        # print(images.shape)
        images = images.reshape(10,28*28)
        outputs_hid256, activations_hid256 = model_hid_256(images.float())

# Reconstruction from testset image
plt.rcParams["figure.figsize"] = (10,6)
i = 6
if i==6:
    fig, (ax1, ax2, ax3, ax4) = plt.subplots(1,4)

    ax1.imshow(images[i].detach().numpy().reshape(28,28), cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs_hid64[i].detach().numpy().reshape(28,28), cmap='gray')
    ax2.set_title('layer 64')
    ax2.axis("off")

    ax3.imshow(outputs_hid128[i].detach().numpy().reshape(28,28), cmap='gray')
    ax3.set_title('layer 128')
    ax3.axis("off")

    ax4.imshow(outputs_hid256[i].detach().numpy().reshape(28,28), cmap='gray')
    ax4.set_title('layer 256')
    ax4.axis("off")

    print("Reconstruction Error in hidden layer 64:", np.dot(((images[i].detach().numpy() -
                                                                ((images[i].detach().numpy() /
                                                                ((images[i].detach().numpy() /

    print("Reconstruction Error in hidden layer 128:", np.dot(((images[i].detach().n
                                                                ((images[i].detach().numpy() /
                                                                ((images[i].detach().numpy() /

    print("Reconstruction Error in hidden layer 256:", np.dot(((images[i].detach().n
                                                                ((images[i].detach().numpy() /
                                                                ((images[i].detach().numpy() /

```


Reconstruction Error in hidden layer 64: 11.198258444172145

Reconstruction Error in hidden layer 128: 7.300481213033256

Reconstruction Error in hidden layer 256: 5.585871585468324

Original Image reconstructed Image-hid layer 64 reconstructed Image-hid layer 128 reconstructed Image-hid layer 256





```
# Reconstruction from non-digit images -- Fashion MNIST o/p
test_data_fashion = torchvision.datasets.FashionMNIST(root=".", train=False, transform=
fashion_image_sample = test_data_fashion.data[9]
```

```
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-images-idx3-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-labels-idx1-ubyte.gz
100%|██████████| 26421880/26421880 [00:01<00:00, 13343461.45it/s]
Extracting ./FashionMNIST/raw/train-images-idx3-ubyte.gz to ./FashionMNIST/raw
```

```
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-images-idx3-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-labels-idx1-ubyte.gz
100%|██████████| 29515/29515 [00:00<00:00, 211148.13it/s]
Extracting ./FashionMNIST/raw/train-labels-idx1-ubyte.gz to ./FashionMNIST/raw
```

```
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-images-idx3-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz
100%|██████████| 4422102/4422102 [00:01<00:00, 3904504.94it/s]
Extracting ./FashionMNIST/raw/t10k-images-idx3-ubyte.gz to ./FashionMNIST/raw
```

```
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-images-idx3-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz
100%|██████████| 5148/5148 [00:00<00:00, 15511693.24it/s]Extracting ./FashionMNIST/raw/t10k-labels-idx1-ubyte.gz to ./FashionMNIST/raw
```

```
model_hid_64.eval()
with torch.no_grad():
    images = fashion_image_sample.reshape(1, 28*28)
    outputs_hid64, _ = model_hid_64(images.float())
```

```
model_hid_128.eval()
with torch.no_grad():
    images = fashion_image_sample.reshape(1, 28*28)
    outputs_hid128, _ = model_hid_128(images.float())
```

```
model_hid_256.eval()
with torch.no_grad():
    images = fashion_image_sample.reshape(1, 28*28)
    outputs_hid256, _ = model_hid_256(images.float())
```

```

outputs_hid256,_ = model_hid_256(images.float())

plt.rcParams["figure.figsize"] = (10,6)

fig, (ax1, ax2, ax3, ax4) = plt.subplots(1,4)


ax1.imshow(fashion_image_sample.detach().numpy().reshape(28,28),cmap='gray')
ax1.set_title('Original Random Image')
ax1.axis("off")

ax2.imshow(outputs_hid64.detach().numpy().reshape(28,28),cmap='gray')
ax2.set_title('layer 64')
ax2.axis("off")

ax3.imshow(outputs_hid128.detach().numpy().reshape(28,28),cmap='gray')
ax3.set_title('layer 128')
ax3.axis("off")

ax4.imshow(outputs_hid256.detach().numpy().reshape(28,28),cmap='gray')
ax4.set_title('layer 256')
ax4.axis("off")
plt.show()

```



```

# random noise images
torch.manual_seed(0)
random_image = torch.randint(low=0, high=255,size=(1,28,28))

model_hid_64.eval()
with torch.no_grad():
    images = random_image.reshape(1,28*28)
    outputs_hid64,_ = model_hid_64(images.float())

model_hid_128.eval()
with torch.no_grad():

```

```

images = random_image.reshape(1,28*28)
outputs_hid128,_ = model_hid_128(images.float())

model_hid_256.eval()
with torch.no_grad():
    images = random_image.reshape(1,28*28)
    outputs_hid256,_ = model_hid_256(images.float())

plt.rcParams["figure.figsize"] = (10,6)

fig, (ax1, ax2, ax3, ax4) = plt.subplots(1,4)


ax1.imshow(random_image.numpy().reshape(28,28),cmap='gray')
ax1.set_title('Ori Random image')
ax1.axis("off")

ax2.imshow(outputs_hid64.detach().numpy().reshape(28,28),cmap='gray')
ax2.set_title('reconst image-hid64')
ax2.axis("off")

ax3.imshow(outputs_hid128.detach().numpy().reshape(28,28),cmap='gray')
ax3.set_title('reconst image-hid128')
ax3.axis("off")

ax4.imshow(outputs_hid256.detach().numpy().reshape(28,28),cmap='gray')
ax4.set_title('reconst image-hid256')
ax4.axis("off")
plt.show()

```



✓ Sparse Autoencoders

over-complete autoencoder (AE) with sparsity regularization using L1 penalty on

```
class Sparse_AE(nn.Module):
    def __init__(self):
        super(Sparse_AE, self).__init__()
        self.encoder = nn.Sequential(
            nn.Linear(784, 1024), # Over-complete layer
            nn.ReLU())
        self.decoder = nn.Sequential(
            nn.Linear(1024, 784), # Reconstruction back to input dimension
            nn.ReLU())

    def forward(self,x):
        x = self.encoder(x) # encoded output
        encoded_output = x
        l1_norm = torch.norm(x, p=1) # calculate L1 norm of the activations
        x = self.decoder(x) # Get reconstructed output
        return x, l1_norm, encoded_output

sparsity_param = 0.5*1e-6 # Sparsity regularization weight
model_sparse1 = Sparse_AE()
criterion_sparse1 = nn.MSELoss()
optimizer_sparse1 = torch.optim.Adam(model_sparse1.parameters(),lr=learning_rate)

training_loss_sparse1 = []

for epoch in range(epochs):
    for images,labels in train_loader:
        images = images.reshape(images.shape[0],-1)
        outputs, l1_norm, _ = model_sparse1(images)
        loss = criterion_sparse1(outputs,images) + sparsity_param * l1_norm
        training_loss_sparse1.append(loss.item())

    optimizer_sparse1.zero_grad()
    loss.backward()
    optimizer_sparse1.step()

    print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.item()))

    Epoch [1/10]: completed, Loss: 0.0129
    Epoch [2/10]: completed, Loss: 0.0102
    Epoch [3/10]: completed, Loss: 0.0091
    Epoch [4/10]: completed, Loss: 0.0086
    Epoch [5/10]: completed, Loss: 0.0076
    Epoch [6/10]: completed, Loss: 0.0079
    Epoch [7/10]: completed, Loss: 0.0071
    Epoch [8/10]: completed, Loss: 0.0075
    Epoch [9/10]: completed, Loss: 0.0065
    Epoch [10/10]: completed, Loss: 0.0070

sparsity_param = 1.5*1e-6
model_sparse2 = Sparse_AE()
```

```
criterion_sparse2 = nn.MSELoss()
optimizer_sparse2 = torch.optim.Adam(model_sparse2.parameters(), lr=learning_rate)

training_loss_sparse2=[]

for epoch in range(epochs):
    for images, labels in train_loader:
        images = images.reshape(images.shape[0], -1)
        outputs, l1_norm, _ = model_sparse2(images)
        loss = criterion_sparse2(outputs, images) + sparsity_param * l1_norm
        training_loss_sparse2.append(loss.item())

    optimizer_sparse2.zero_grad()
    loss.backward()
    optimizer_sparse2.step()

print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.item()))

Epoch [1/10]: completed, Loss: 0.0171
Epoch [2/10]: completed, Loss: 0.0144
Epoch [3/10]: completed, Loss: 0.0123
Epoch [4/10]: completed, Loss: 0.0114
Epoch [5/10]: completed, Loss: 0.0110
Epoch [6/10]: completed, Loss: 0.0106
Epoch [7/10]: completed, Loss: 0.0093
Epoch [8/10]: completed, Loss: 0.0092
Epoch [9/10]: completed, Loss: 0.0088
Epoch [10/10]: completed, Loss: 0.0085
```

```
sparsity_param = 3*1e-6 # Sparsity regularization weight
model_sparse3 = Sparse_AE()
criterion_sparse3 = nn.MSELoss()
optimizer_sparse3 = torch.optim.Adam(model_sparse3.parameters(), lr=learning_rate)

training_loss_sparse3 = []

for epoch in range(epochs):
    for images, labels in train_loader:
        images = images.reshape(images.shape[0], -1)
        outputs, l1_norm, _ = model_sparse3(images)
        loss = criterion_sparse3(outputs, images) + sparsity_param * l1_norm
        training_loss_sparse3.append(loss.item())

    optimizer_sparse3.zero_grad()
    loss.backward()
    optimizer_sparse3.step()


print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.item()))

Epoch [1/10]: completed, Loss: 0.0249
Epoch [2/10]: completed, Loss: 0.0196
```



```
Epoch [3/10]: completed, Loss: 0.0187
Epoch [4/10]: completed, Loss: 0.0174
Epoch [5/10]: completed, Loss: 0.0162
Epoch [6/10]: completed, Loss: 0.0151
Epoch [7/10]: completed, Loss: 0.0142
Epoch [8/10]: completed, Loss: 0.0139
Epoch [9/10]: completed, Loss: 0.0142
Epoch [10/10]: completed, Loss: 0.0132
```

```
plt.plot(range(1, len(training_loss_sparse3)+1), training_loss_sparse1, label="lambda=0.5*1e-6")
plt.plot(range(1, len(training_loss_sparse3)+1), training_loss_sparse2, label="lambda=1.5*1e-6")
plt.plot(range(1, len(training_loss_sparse3)+1), training_loss_sparse3, label="lambda=3*1e-6")
plt.legend()
plt.grid()
plt.title("Training Loss Vs Iterations")
plt.xlabel("Number of Iterations")
plt.ylabel("Training Loss")
plt.show()
```



```
model_sparse1.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        outputs_sparse1, _ , activation_sparse1 = model_sparse1(images.float())
```

```
model_sparse2.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        outputs_sparse2, _ , activation_sparse2 = model_sparse2(images.float())
```

```
model_sparse3.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        outputs_sparse3, _ , activation_sparse3 = model_sparse3(images.float())
```

```
plt.rcParams["figure.figsize"] = (10,6)
```

```
for i in range(5):
    fig, (ax1, ax2, ax3, ax4) = plt.subplots(1,4)

    ax1.imshow(images[i].detach().numpy().reshape(28,28), cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs_sparse1[i].detach().numpy().reshape(28,28), cmap='gray')
    ax2.set_title('lambda=0.5*1e-6')
    ax2.axis("off")

    ax3.imshow(outputs_sparse2[i].detach().numpy().reshape(28,28), cmap='gray')
    ax3.set_title('lambda=1.5*1e-6')
    ax3.axis("off")

    ax4.imshow(outputs_sparse3[i].detach().numpy().reshape(28,28), cmap='gray')
    ax4.set_title('lambda=3*1e-6')
    ax4.axis("off")
```


```
plt.show()
```

```
print("Reconstruction Error in SparseAE lambda = 0.6*1e-6 :",np.dot(((images[i].c
                                                                    ((images[i].detac
```


```
print("Reconstruction Error in SparseAE lambda = 1.5*1e-6:",np.dot(((images[i].de
                                                                    ((images[i].det
```

```
print("Reconstruction Error in SparseAE lambda = 3*1e-6:", np.dot(((images[i].data
                                                                    ((images[i].data
```


Original Image




lambda=0.5*1e-6



lambda=1.5*1e-6



lambda=3*1e-6



Reconstruction Error in SparseAE lambda = 0.6*1e-6 : 2.208584472857042


Reconstruction Error in SparseAE lambda = 1.5*1e-6: 1.9335901925147088

Reconstruction Error in SparseAE lambda = 3*1e-6: 3.6661419552609127


Original Image




lambda=0.5*1e-6



lambda=1.5*1e-6



lambda=3*1e-6



Reconstruction Error in SparseAE lambda = 0.6*1e-6 : 7.516961848370433

Reconstruction Error in SparseAE lambda = 1.5*1e-6: 7.161093606775397

Reconstruction Error in SparseAE lambda = 3*1e-6: 11.231844113122616


Original Image



lambda=0.5*1e-6



lambda=1.5*1e-6



lambda=3*1e-6




Reconstruction Error in SparseAE lambda = 0.6*1e-6 : 11.913395591403894


Reconstruction Error in SparseAE lambda = 1.5*1e-6: 9.469317564296931

Reconstruction Error in SparseAE lambda = 3*1e-6: 9.207602082149549

Original Image



lambda=0.5*1e-6



lambda=1.5*1e-6



lambda=3*1e-6




Reconstruction Error in SparseAE lambda = 0.6*1e-6 : 5.362101635436673

Reconstruction Error in SparseAE lambda = 1.5*1e-6: 4.511604371463333314631


Reconstruction Error in SparseAE lambda = 1.5×10^{-6} : 4.371462238244634

Reconstruction Error in SparseAE lambda = 3×10^{-6} : 7.5787287819801605

Original Image




lambda= 0.5×10^{-6}



lambda= 1.5×10^{-6}



lambda= 3×10^{-6}



Reconstruction Error in SparseAE lambda = 0.6×10^{-6} : 4.129679657692563

Reconstruction Error in SparseAE lambda = 1.5×10^{-6} : 5.395042252176756

Reconstruction Error in SparseAE lambda = 3×10^{-6} : 12.234988018390908

```
# visualize the learned filters of the Sparse AE as images
```

```
plt.rcParams["figure.figsize"] = (15,6)
```

```
for i in range(5):
```

```
    fig, (ax1, ax2, ax3, ax4, ax5) = plt.subplots(1,5)
```

```
    ax1.imshow(images[i].detach().numpy().reshape(28,28), cmap='gray')
```

```
    ax1.set_title('Original Image')
```

```
    ax1.axis("off")
```

```
    ax2.imshow(np.array(activations_hid256.detach().numpy())[i].reshape(int(np.sqrt
```

```
    ax2.set_title('Vanialla AE')
```

```
    ax2.axis("off")
```

```
    ax3.imshow(np.array(activation_sparse1.detach().numpy())[i].reshape(int(np.sqrt
```

```
    ax3.set_title('lambda= $0.6 \times 10^{-6}$ ')
```

```
    ax3.axis("off")
```

```
    ax4.imshow(np.array(activation_sparse2.detach().numpy())[i].reshape(int(np.sqrt
```

```
    ax4.set_title('lambda= $1.5 \times 10^{-6}$ ')
```

```
    ax4.axis("off")
```


```
    ax5.imshow(np.array(activation_sparse3.detach().numpy())[i].reshape(int(np.sqrt
```

```
    ax5.set_title('lambda= $3 \times 10^{-6}$ ')
```

```
    ax5.axis("off")
```

```
fig.tight_layout()
```

```
plt.show()
```



✓ Denoising Autoencoders

```
class Denoising_AE(nn.Module):
    def __init__(self):
        super(Denoising_AE, self).__init__()
        self.encoder = nn.Sequential( # encoder
            nn.Linear(784, 256),
            nn.ReLU())
        self.decoder = nn.Sequential( # decoder
            nn.Linear(256, 784),
            nn.ReLU())

    def forward(self,x):
        x = self.encoder(x)
        x = self.decoder(x)
        return x

def add_noise(img, noise_val):
    noise = torch.randn(img.size()) * noise_val
    noisy_img = img + noise
    return noisy_img

# Hyperparameters
learning_rate = 0.0001
epochs = 10
batch_size = 200

model_denoise1 = Denoising_AE()
criterion_denoise1 = nn.MSELoss()
```

```

criterion_denoise1 = nn.MSELoss()
optimizer_denoise1 = torch.optim.Adam(model_denoise1.parameters(), lr=learning_rate)

training_loss_denoise1 = []

for epoch in range(epochs):
    for images, labels in train_loader:
        images = images.reshape(images.shape[0], -1)
        noisy_images = add_noise(images, 0.3)
        outputs = model_denoise1(noisy_images)
        loss = criterion_denoise1(outputs, images)
        training_loss_denoise1.append(loss.item())

    optimizer_denoise1.zero_grad()
    loss.backward()
    optimizer_denoise1.step()

    print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.item()))

```

```

Epoch [1/10]: completed, Loss: 0.0285
Epoch [2/10]: completed, Loss: 0.0188
Epoch [3/10]: completed, Loss: 0.0142
Epoch [4/10]: completed, Loss: 0.0137
Epoch [5/10]: completed, Loss: 0.0121
Epoch [6/10]: completed, Loss: 0.0116
Epoch [7/10]: completed, Loss: 0.0114
Epoch [8/10]: completed, Loss: 0.0113
Epoch [9/10]: completed, Loss: 0.0107
Epoch [10/10]: completed, Loss: 0.0108

```

```

model_denoise2 = Denoising_AE()
criterion_denoise2 = nn.MSELoss()
optimizer_denoise2 = torch.optim.Adam(model_denoise2.parameters(), lr=learning_rate)

training_loss_denoise2 = []

for epoch in range(epochs):
    for images, labels in train_loader:
        images = images.reshape(images.shape[0], -1)
        noisy_images = add_noise(images, 0.5)
        outputs = model_denoise2(noisy_images)
        loss = criterion_denoise2(outputs, images)
        training_loss_denoise2.append(loss.item())

    optimizer_denoise2.zero_grad()
    loss.backward()
    optimizer_denoise2.step()

    print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.item()))

```

```

Epoch [1/10]: completed, Loss: 0.0349
Epoch [2/10]: completed, Loss: 0.0257
Epoch [3/10]: completed, Loss: 0.0215

```

```
Epoch [9/10]: completed, Loss: 0.0177
Epoch [4/10]: completed, Loss: 0.0194
Epoch [5/10]: completed, Loss: 0.0182
Epoch [6/10]: completed, Loss: 0.0187
Epoch [7/10]: completed, Loss: 0.0176
Epoch [8/10]: completed, Loss: 0.0172
Epoch [9/10]: completed, Loss: 0.0177
Epoch [10/10]: completed, Loss: 0.0178
```

```
model_denoise3 = Denoising_AE()
criterion_denoise3 = nn.MSELoss()
optimizer_denoise3 = torch.optim.Adam(model_denoise3.parameters(),lr=learning_rate)
```

```
training_loss_denoise3 = []
```

```
for epoch in range(epochs):
    for images,labels in train_loader:
        images = images.reshape(images.shape[0],-1)
        noisy_images = add_noise(images, 0.8)
        outputs = model_denoise3(noisy_images)
        loss = criterion_denoise3(outputs, images)
        training_loss_denoise3.append(loss.item())

    optimizer_denoise3.zero_grad()
    loss.backward()
    optimizer_denoise3.step()
    print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.item()))
```

```
Epoch [1/10]: completed, Loss: 0.0437
Epoch [2/10]: completed, Loss: 0.0355
Epoch [3/10]: completed, Loss: 0.0312
Epoch [4/10]: completed, Loss: 0.0293
Epoch [5/10]: completed, Loss: 0.0276
Epoch [6/10]: completed, Loss: 0.0280
Epoch [7/10]: completed, Loss: 0.0281
Epoch [8/10]: completed, Loss: 0.0266
Epoch [9/10]: completed, Loss: 0.0260
Epoch [10/10]: completed, Loss: 0.0238
```

```
model_denoise4 = Denoising_AE()
criterion_denoise4 = nn.MSELoss()
optimizer_denoise4 = torch.optim.Adam(model_denoise4.parameters(),lr=learning_rate)
```

```
training_loss_denoise4 = []
```

```
for epoch in range(epochs):
    for images,labels in train_loader:
        images = images.reshape(images.shape[0],-1)
        noisy_images = add_noise(images, 0.9)
        outputs = model_denoise4(noisy_images)
        loss = criterion_denoise4(outputs, images)
```



```
training_loss_denoise4.append(loss.item())


optimizer_denoise4.zero_grad()
loss.backward()
optimizer_denoise4.step()
print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.item()))

Epoch [1/10]: completed, Loss: 0.0513
Epoch [2/10]: completed, Loss: 0.0390
Epoch [3/10]: completed, Loss: 0.0348
Epoch [4/10]: completed, Loss: 0.0324
Epoch [5/10]: completed, Loss: 0.0317
Epoch [6/10]: completed, Loss: 0.0287
Epoch [7/10]: completed, Loss: 0.0285
Epoch [8/10]: completed, Loss: 0.0289
Epoch [9/10]: completed, Loss: 0.0288
Epoch [10/10]: completed, Loss: 0.0276
```

```
plt.rcParams["figure.figsize"] = (10, 6)
```

```
plt.plot(range(1, len(training_loss_denoise1)+1), training_loss_denoise1, label="noise factor = 0.3")
plt.plot(range(1, len(training_loss_denoise1)+1), training_loss_denoise2, label="noise factor = 0.5")
plt.plot(range(1, len(training_loss_denoise1)+1), training_loss_denoise3, label="noise factor = 0.8")
plt.plot(range(1, len(training_loss_denoise1)+1), training_loss_denoise4, label="noise factor = 0.9")
plt.legend()
plt.grid()
plt.title("Training Loss Vs Iterations")
plt.xlabel("No: of Iterations")
plt.ylabel("Training Loss")
plt.show()
```





```
# With different noise factor 0.3,0.5,0.8,0.9 given to Vanilla AE
model_hid_256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images = add_noise(images,0.3)
        outputs_hid256_01, activations_hid256 = model_hid_256(noisy_images.float())

model_hid_256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images = add_noise(images,0.5)
        outputs_hid256_02, activations_hid256 = model_hid_256(noisy_images.float())

model_hid_256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images = add_noise(images, 0.8)
        outputs_hid256_03, activations_hid256 = model_hid_256(noisy_images.float())

model_hid_256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images = add_noise(images, 0.9)
        outputs_hid256_04, activations_hid256 = model_hid_256(noisy_images.float())

plt.rcParams["figure.figsize"] = (10,6)
i = 7
fig, (ax1, ax2,ax3,ax4,ax5) = plt.subplots(1,5)
```

```

ax1.imshow(images[i].detach().numpy().reshape(28,28), cmap='gray')
ax1.set_title('Original Image')
ax1.axis("off")

ax2.imshow(outputs_hid256_01[i].detach().numpy().reshape(28,28), cmap='gray')
ax2.set_title('noise factor=0.3')
ax2.axis("off")

ax3.imshow(outputs_hid256_02[i].detach().numpy().reshape(28,28), cmap='gray')
ax3.set_title('noise factor=0.5')
ax3.axis("off")

ax4.imshow(outputs_hid256_03[i].detach().numpy().reshape(28,28), cmap='gray')
ax4.set_title('noise factor=0.8')
ax4.axis("off")

ax5.imshow(outputs_hid256_04[i].detach().numpy().reshape(28,28), cmap='gray')
ax5.set_title('noise factor=0.9')
ax5.axis("off")

print("Reconstruction Error in VanillaAE with noise factor = 0.3 :", np.dot(((image
                                                                    ((image

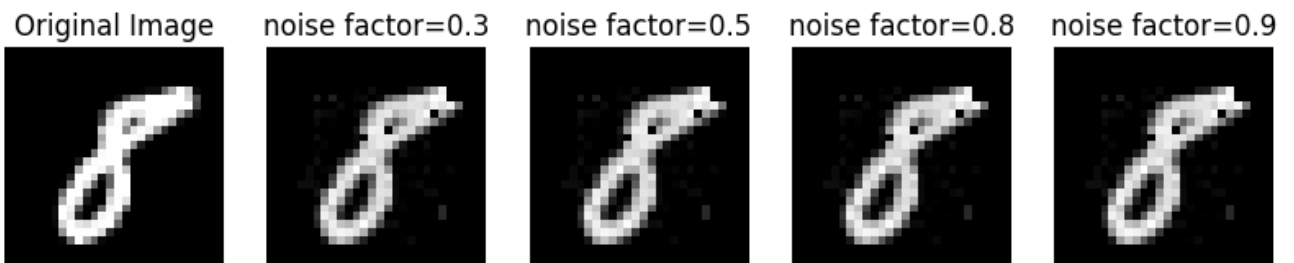
print("Reconstruction Error in VanillaAE with noise factor = 0.5 :", np.dot(((image
                                                                    ((image

print("Reconstruction Error in VanillaAE with noise factor = 0.8 :", np.dot(((image
                                                                    ((image

print("Reconstruction Error in VanillaAE with noise factor = 0.9 :", np.dot(((image
                                                                    ((image

Reconstruction Error in VanillaAE with noise factor = 0.3 : 4.758285692172574
Reconstruction Error in VanillaAE with noise factor = 0.5 : 4.770937421211641!
Reconstruction Error in VanillaAE with noise factor = 0.8 : 4.758050556874504
Reconstruction Error in VanillaAE with noise factor = 0.9 : 4.75845746901802

```




2. Manifold learning

Manifold Learning

```
# High-dimensional spaces are sparse: In a 784-dimensional space, valid digit ima  
# Random noise disrupts correlations between pixels  
# Perception in high dimensions: In high-dimensional spaces, distances between po
```

```
# Select a random digit  
random_digit, _ = next(iter(train_loader))  
  
# Original digit  
original_digit = random_digit[0].squeeze().numpy() # Shape (28, 28)  
  
# Add random noise  
noise_factor = 0.5  
random_noise = np.random.normal(loc=0.0, scale=noise_factor, size=original_digit.  
noisy_digit = np.clip(original_digit + random_noise, 0.0, 1.0) # Keep values in  
  
# Plotting the original and noisy digits  
plt.figure(figsize=(6, 3))  
  
# Original digit  
plt.subplot(1, 2, 1)  
plt.imshow(original_digit, cmap='gray')  
plt.title('Original Digit')  
plt.axis('off')  
  
# Noisy digit  
plt.subplot(1, 2, 2)  
plt.imshow(noisy_digit, cmap='gray')  
plt.title('Noisy Digit')  
plt.axis('off')  
  
plt.show()
```



```

class AE_manifold(nn.Module):
    def __init__(self):
        super(AE_manifold, self).__init__()
        self.encoder = nn.Sequential(
            nn.Linear(784, 64),
            nn.ReLU(),
            nn.Linear(64, 8)
        )
        self.decoder = nn.Sequential(
            nn.Linear(8, 64),
            nn.ReLU(),
            nn.Linear(64, 784),
            nn.Sigmoid() # Sigmoid to ensure output is in range [0, 1]
        )

    def forward(self, x):
        latent_representation = self.encoder(x) # Get latent representation
        reconstructed_image = self.decoder(latent_representation) # Reconstruct
        return reconstructed_image, latent_representation # Return both reconstr

# Initialize model, loss function, and optimizer
model = AE_manifold()
criterion = nn.MSELoss()
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

# Training loop
training_loss_manifold = []
for epoch in range(epochs):
    for images, labels in train_loader:
        images = images.reshape(images.shape[0], -1) # 2D to 1D vector
        outputs, _ = model(images) # Reconstructed images and latent representat
        loss = criterion(outputs, images) # Reconstruction error
        training_loss_manifold.append(loss.item())

        optimizer.zero_grad() # Clears the gradients from the previous batch
        loss.backward() # Backpropagation
        optimizer.step() # Update the model parameters

print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.i

Epoch [1/10]: completed, Loss: 0.0731
Epoch [2/10]: completed, Loss: 0.0668
Epoch [3/10]: completed, Loss: 0.0650
Epoch [4/10]: completed, Loss: 0.0565
Epoch [5/10]: completed, Loss: 0.0487
Epoch [6/10]: completed, Loss: 0.0465

```

```
Epoch [7/10]: completed, Loss: 0.0438
Epoch [8/10]: completed, Loss: 0.0428
Epoch [9/10]: completed, Loss: 0.0397
Epoch [10/10]: completed, Loss: 0.0378
```

```
# # Autoencoder class
# class AE_manifold(nn.Module):
#     def __init__(self):
#         super(AE_manifold, self).__init__()
#         self.encoder = nn.Sequential(
#             nn.Linear(784, 64),
#             nn.ReLU(),
#             nn.Linear(64, 8)
#         )
#         self.decoder = nn.Sequential(
#             nn.Linear(8, 64),
#             nn.ReLU(),
#             nn.Linear(64, 784),
#             nn.Sigmoid() # Sigmoid to ensure output is in range [0, 1]
#         )

#     def forward(self, x):
#         x = self.encoder(x)
#         x = self.decoder(x)
#         return x

# # Initialize model, loss function, and optimizer
# model = AE_manifold()
# criterion = nn.MSELoss()
# optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

# # Training loop
# training_loss_manifold = []
# for epoch in range(epochs):
#     for images, labels in train_loader:
#         images = images.reshape(images.shape[0], -1) # 2D to 1D vector
#         outputs = model(images) # Reconstructed images (only one value returned)
#         loss = criterion(outputs, images) # Reconstruction error
#         training_loss_manifold.append(loss.item())

#         optimizer.zero_grad() # Clears the gradients from the previous batch
#         loss.backward() # Backpropagation
#         optimizer.step() # Update the model parameters

#     print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss))

# # Testing on the test set
# test_image = test_dataset[0][0].unsqueeze(0) # Take a test image and add a batch
```

```

test_image = test_dataset[19/65].unsqueeze(0) # Take a test image and add a batch
test_image = test_image.float() / 255.0 # Normalize

# Pass through encoder to get latent representation
with torch.no_grad():
    reconstructed_image, latent_representation = model(test_image) # Now it shou

# Add noise to the latent representation
noise_factor = 0.5
noisy_latent = latent_representation + torch.randn(latent_representation.shape) *

# Reconstruct the image from the noisy latent representation
noisy_reconstructed_image = model.decoder(noisy_latent)


# Reshape images for plotting (28x28)
original_image = test_image.view(28, 28).detach().numpy()
noisy_image = noisy_reconstructed_image.view(28, 28).detach().numpy()

# Plot the original and noisy reconstructed images
plt.figure(figsize=(6, 3))
plt.subplot(1, 2, 1)
plt.imshow(original_image, cmap='gray')
plt.title('Original Image')
plt.axis('off')

plt.subplot(1, 2, 2)
plt.imshow(noisy_image, cmap='gray')
plt.title('Noisy Reconstructed Image')
plt.axis('off')

plt.show()

```



✓ Convolutional Autoencoders

Convolutional Autoencoders

```

class ConvolutionAE_unpooling(nn.Module):  # class definition
    def __init__(self):  # constructor
        super(ConvolutionAE_unpooling, self).__init__()

        # encoder module
        self.enc_conv1 = nn.Sequential(
            nn.Conv2d(1, 8, kernel_size = 3, stride = 1, padding = 1), # 28x28x1 to 14x14x8
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2, 2), return_indices = True)
        )
        self.enc_conv2 = nn.Sequential(
            nn.Conv2d(8, 16, kernel_size = 3, stride = 1, padding = 1), # 14x14x8 to 7x7x16
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2, 2), return_indices = True)
        )
        self.enc_conv3 = nn.Sequential(
            nn.Conv2d(16, 16, kernel_size = 3, stride = 1, padding = 1), # 7x7x16 to 3x3x16
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2, 2), return_indices = True)
        )

        # decoder module
        self.dec_conv1 = nn.Sequential(nn.Identity()) # 7x7x16 to 7x7x16
        self.dec_conv2 = nn.Sequential(
            nn.Conv2d(16, 8, kernel_size = 3, stride = 1, padding = 1), # 14x14x16 to 7x7x8
            nn.ReLU()
        )
        self.dec_conv3 = nn.Sequential(
            nn.Conv2d(8, 1, kernel_size = 3, stride = 1, padding = 1), # 28x28x8 to 28x28x1
            nn.ReLU()
        )

        # unpooling operation
        self.unpool = nn.MaxUnpool2d(kernel_size = (2, 2))

    def forward(self, x):  # forward pass
        enc_input, indices1 = self.enc_conv1(x.float()) # 28x28x1 to 14x14x8
        enc_input, indices2 = self.enc_conv2(enc_input) # 14x14x8 to 7x7x16
        enc_input, indices3 = self.enc_conv3(enc_input) # 7x7x16 to 3x3x16

        reconst_img = self.unpool(enc_input, indices3, output_size=torch.Size([batch_size, 1, 28, 28]))
        reconst_img = self.dec_conv1(reconst_img) # 7x7x16 to 7x7x16
        reconst_img = self.unpool(reconst_img, indices2) # 7x7x16 to 14x14x16
        reconst_img = self.dec_conv2(reconst_img) # 14x14x16 to 14x14x8
        reconst_img = self.unpool(reconst_img, indices1) # 14x14x8 to 28x28x8
        reconst_img = self.dec_conv3(reconst_img) # 28x28x8 to 28x28x1

```



```
return reconst_img, enc_input
```

```
model_conAE1 = ConvolutionAE_unpooling()
criterion_conAE1 = nn.MSELoss()
optimizer_conAE1 = torch.optim.Adam(model_conAE1.parameters(), lr=0.001)
```

```
training_loss_conAE1 = []
```

```
for epoch in range(epochs):
    for images, labels in train_loader:
        outputs, _ = model_conAE1(images)
        loss = criterion_conAE1(outputs, images)
        training_loss_conAE1.append(loss.item())

        optimizer_conAE1.zero_grad()
        loss.backward()
        optimizer_conAE1.step()
```

```
print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.item()))
```

```
Epoch [1/10]: completed, Loss: 0.0146
Epoch [2/10]: completed, Loss: 0.0103
Epoch [3/10]: completed, Loss: 0.0091
Epoch [4/10]: completed, Loss: 0.0079
Epoch [5/10]: completed, Loss: 0.0068
Epoch [6/10]: completed, Loss: 0.0069
Epoch [7/10]: completed, Loss: 0.0059
Epoch [8/10]: completed, Loss: 0.0061
Epoch [9/10]: completed, Loss: 0.0057
Epoch [10/10]: completed, Loss: 0.0055
```

```
class ConvolutionAE_deconv(nn.Module):
    def __init__(self):
        super(ConvolutionAE_deconv, self).__init__()

        # encoder
        self.enc_conv1 = nn.Sequential(
            nn.Conv2d(1, 8, kernel_size = 3, stride = 1, padding = 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2, 2))
        )
        self.enc_conv2 = nn.Sequential(
            nn.Conv2d(8, 16, kernel_size = 3, stride = 1, padding = 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2, 2))
        )
        self.enc_conv3 = nn.Sequential(
            nn.Conv2d(16, 16, kernel_size = 3, stride = 1, padding = 1),
            nn.ReLU()
```

```

        nn.ReLU(),
        nn.MaxPool2d(kernel_size = (2,2))
    )

# decoder
self.dec_conv1 = nn.Sequential(
    nn.ConvTranspose2d(16,16, kernel_size = 3, stride = 2),
    nn.ReLU()
)
self.dec_conv2 = nn.Sequential(
    nn.ConvTranspose2d(16,8, kernel_size = 4, stride = 2, padding = 1),
    nn.ReLU()
)
self.dec_conv3 = nn.Sequential(
    nn.ConvTranspose2d(8,1, kernel_size = 4, stride = 2, padding = 1),
    nn.ReLU()
)

def forward(self,x):
    enc_input = self.enc_conv1(x.float())
    enc_input = self.enc_conv2(enc_input)
    enc_input = self.enc_conv3(enc_input)

    reconst_img = self.dec_conv1(enc_input)
    reconst_img = self.dec_conv2(reconst_img)
    reconst_img = self.dec_conv3(reconst_img)

    return reconst_img, enc_input

model_conAE2 = ConvolutionAE_deconv()
criterion_conAE2 = nn.MSELoss()
optimizer_conAE2 = torch.optim.Adam(model_conAE2.parameters(),lr=0.001)

training_loss_conAE2=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        outputs, _ = model_conAE2(images)
        loss = criterion_conAE2(outputs, images)
        training_loss_conAE2.append(loss.item())

        optimizer_conAE2.zero_grad()
        loss.backward()
        optimizer_conAE2.step()

    print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.item()))

Epoch [1/10]: completed, Loss: 0.0245
Epoch [2/10]: completed, Loss: 0.0183

```

```
Epoch [2/10]: completed, Loss: 0.0165  
Epoch [3/10]: completed, Loss: 0.0164  
Epoch [4/10]: completed, Loss: 0.0151  
Epoch [5/10]: completed, Loss: 0.0134  
Epoch [6/10]: completed, Loss: 0.0124  
Epoch [7/10]: completed, Loss: 0.0129  
Epoch [8/10]: completed, Loss: 0.0119  
Epoch [9/10]: completed, Loss: 0.0120  
Epoch [10/10]: completed, Loss: 0.0117
```

```
class ConvolutionAE_unpool_deconv(nn.Module):  
    def __init__(self):  
        super(ConvolutionAE_unpool_deconv,self).__init__()  
  
        # encoder  
        self.enc_conv1 = nn.Sequential(  
            nn.Conv2d(1,8, kernel_size = 3, stride = 1, padding = 1),  
            nn.ReLU(),  
            nn.MaxPool2d(kernel_size = (2,2),return_indices = True)  
        )  
        self.enc_conv2 = nn.Sequential(  
            nn.Conv2d(8,16, kernel_size = 3, stride = 1, padding = 1),  
            nn.ReLU(),  
            nn.MaxPool2d(kernel_size = (2,2),return_indices = True)  
        )  
        self.enc_conv3 = nn.Sequential(  
            nn.Conv2d(16,16, kernel_size = 3, stride = 1, padding = 1),  
            nn.ReLU(),nn.MaxPool2d(kernel_size = (2,2),return_indices = True)  
        )  
  
        # decoder  
        self.dec_conv1 = nn.Sequential(  
            nn.ConvTranspose2d(16,16, kernel_size = 3, stride = 1, padding = 1),  
            nn.ReLU()  
        )  
        self.dec_conv2 = nn.Sequential(  
            nn.ConvTranspose2d(16,8, kernel_size = 3, stride = 1, padding = 1),  
            nn.ReLU()  
        )  
        self.dec_conv3 = nn.Sequential(  
            nn.ConvTranspose2d(8,1, kernel_size = 3, stride = 1, padding = 1),  
            nn.ReLU()  
        )  
        # unpooling  
        self.unpool = nn.MaxUnpool2d(kernel_size = (2,2))  
  
    def forward(self,x):  
        enc_input,indices1 = self.enc_conv1(x.float())  
        enc_input,indices2 = self.enc_conv2(enc_input)  
        enc_input,indices3 = self.enc_conv3(enc_input)
```

```

reconst_img = self.unpool(enc_input, indices3, output_size=torch.Size([batch_s
reconst_img = self.dec_conv1(reconst_img )
reconst_img = self.unpool(reconst_img , indices2)
reconst_img = self.dec_conv2(reconst_img )
reconst_img = self.unpool(reconst_img , indices1)
reconst_img = self.dec_conv3(reconst_img )

```

```

return reconst_img, enc_input

```

```

model_conAE3 = ConvolutionAE_unpool_deconv()
criterion_conAE3 = nn.MSELoss()
optimizer_conAE3 = torch.optim.Adam(model_conAE3.parameters(), lr=0.001)

```

```

training_loss_conAE3 = []

```

```

for epoch in range(epochs):
    for images, labels in train_loader:
        outputs, _ = model_conAE3(images)
        loss = criterion_conAE3(outputs, images)
        training_loss_conAE3.append(loss.item())

```

```

optimizer_conAE3.zero_grad()
loss.backward()
optimizer_conAE3.step()

```

```

print("Epoch [{}/{}]: completed, Loss: {:.4f}".format(epoch+1, epochs, loss.item()))

```

```

Epoch [1/10]: completed, Loss: 0.0150
Epoch [2/10]: completed, Loss: 0.0105
Epoch [3/10]: completed, Loss: 0.0085
Epoch [4/10]: completed, Loss: 0.0078
Epoch [5/10]: completed, Loss: 0.0065
Epoch [6/10]: completed, Loss: 0.0060
Epoch [7/10]: completed, Loss: 0.0053
Epoch [8/10]: completed, Loss: 0.0047
Epoch [9/10]: completed, Loss: 0.0041
Epoch [10/10]: completed, Loss: 0.0039


```

```

plt.plot(range(1, len(training_loss_conAE1)+1), training_loss_conAE1, label="ConvAE1")
plt.plot(range(1, len(training_loss_conAE2)+1), training_loss_conAE2, label="ConvAE2")
plt.plot(range(1, len(training_loss_conAE3)+1), training_loss_conAE3, label="ConvAE3")
plt.legend()
plt.grid()
plt.title("Training Loss Vs Iterations")
plt.xlabel("No. of Iterations")
plt.ylabel("Training Loss")
plt.show()

```

Training Loss Vs Iterations



```
model_conAE1.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,1,28,28)
        outputs_conAE1, _ = model_conAE1(images.float())
```

```
model_conAE2.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,1,28,28)
        outputs_conAE2, _ = model_conAE2(images.float())
```

```
activation_conAE3 = []
model_conAE3.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,1,28,28)
        outputs_conAE3, _ = model_conAE3(images.float())
```

```

plt.rcParams["figure.figsize"] = (10,6)

for i in range(5):
    fig, (ax1, ax2, ax3, ax4) = plt.subplots(1,4)

    ax1.imshow(images[i].detach().numpy().reshape(28,28), cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs_conAE1[i].detach().numpy().reshape(28,28), cmap='gray')
    ax2.set_title('CAE_Unpooling')
    ax2.axis("off")

    ax3.imshow(outputs_conAE2[i].detach().numpy().reshape(28,28), cmap='gray')
    ax3.set_title('CAE_Deconv')
    ax3.axis("off")

    ax4.imshow(outputs_conAE3[i].detach().numpy().reshape(28,28), cmap='gray')
    ax4.set_title('CAE_Unpooling+Deconv')
    ax4.axis("off")


plt.show()

print("Reconstruction Error in ConvAE Unpooling:", np.sum(np.dot(((images[i].det
((images[i].detach().nu


print("Reconstruction Error in ConvAE Deconv:", np.sum(np.dot(((images[i].detach
((images[i].detach().n

print("Reconstruction Error in ConvAE Unpooling+Deconv:", np.sum(np.dot(((images
((images[i].detac

```




Reconstruction Error in ConvAE Unpooling: 41.09798739985503
 Reconstruction Error in ConvAE Deconv: 11.099846611587598
 Reconstruction Error in ConvAE Unpooling+Deconv: 157.03552962947595







Reconstruction Error in ConvAE Unpooling: 249.00256747616373
 Reconstruction Error in ConvAE Deconv: 49.661515593520456
 Reconstruction Error in ConvAE Unpooling+Deconv: 906.3791487199521



Reconstruction Error in ConvAE Unpooling: 236.94265814368208
 Reconstruction Error in ConvAE Deconv: 34.88963381365123
 Reconstruction Error in ConvAE Unpooling+Deconv: 708.2955451104583



Reconstruction Error in ConvAE Unpooling: 191.92509100020084
 Reconstruction Error in ConvAE Deconv: 7.73961747405001
 Reconstruction Error in ConvAE Unpooling+Deconv: 538.5825585477866



Reconstruction Error in ConvAE Unpooling: 79.80316303619999
 Reconstruction Error in ConvAE Deconv: -4.719208042428017
 Reconstruction Error in ConvAE Unpooling+Deconv: 256.2937023179323

```
model_conAE1.enc_conv1[0].weight.detach().numpy().squeeze().shape
```

```
(8, 3, 3)
```

```
# Function for visualisation of weights
```

```
def visualize_tensor(tensor, ch=0, allkernels=False, nrow=8, padding=1):
```

```
    n,c,w,h = tensor.shape
```

```
    if allkernels: tensor = tensor.view(n*c, -1, w, h)
```

```
    elif c != 3: tensor = tensor[:,ch,:,:].unsqueeze(dim=1)
```

```
    rows = np.min((tensor.shape[0] // nrow + 1, 64))
```

```
    grid = utils.make_grid(tensor, nrow=nrow, normalize=True, padding=padding)
```

```
    plt.figure( figsize=(nrow,rows) )
```

```
    plt.imshow(grid.numpy().transpose((1, 2, 0)))
```

```
# Visualising decoder weights for convolution Autoencoder with unpooling
```

```
filter = model_conAE1.dec_conv2[0].weight.data.clone()
```

```
visualize_tensor(filter, ch=0, allkernels=False)
```

```
plt.axis('off')
```

```
plt.ioff()
```

```
plt.title('decoder_conv2 Weights')
```

```
plt.show()
```

```
filter = model_conAE1.dec_conv3[0].weight.data.clone()
```

```
visualize_tensor(filter, ch=0, allkernels=False)
```


```
plt.axis('off')
```

```
plt.ioff()
```

```
plt.title('decoder_conv3 Weights')
```

```
plt.show()
```

decoder_conv2 Weights



decoder_conv3 Weights





```
# Visualising decoder weights for convolution Autoencoders with deconvolution
```

```
filter = model_conAE2.dec_conv1[0].weight.data.clone()
visualize_tensor(filter, ch=0, allkernels=False)
```

```
plt.axis('off')
plt.ioff()
plt.title('decoder_conv1 Weights')
plt.show()
```


```
filter = model_conAE2.dec_conv2[0].weight.data.clone()
visualize_tensor(filter, ch=0, allkernels=False)
```

```
plt.axis('off')
plt.ioff()
plt.title('decoder_conv2 Weights')
plt.show()
```

```
filter = model_conAE2.dec_conv3[0].weight.data.clone()
visualize_tensor(filter, ch=0, allkernels=False)
```


```
plt.axis('off')
plt.ioff()
plt.title('decoder_conv3 Weights')
plt.show()
```

decoder_conv1 Weights



decoder_conv2 Weights





```
# Visualising decoder weights for convolution Autoencoders with unpooling + deconv
```

```
filter = model_conAE3.dec_conv1[0].weight.data.clone()
visualize_tensor(filter, ch=0, allkernels=False)
```


```
plt.axis('off')
plt.ioff()
plt.title('decoder_conv1 Weights')
plt.show()
```

```
filter = model_conAE3.dec_conv2[0].weight.data.clone()
visualize_tensor(filter, ch=0, allkernels=False)
```


```
plt.axis('off')
plt.ioff()
plt.title('decoder_conv2 Weights')
plt.show()
```

```
filter = model_conAE3.dec_conv3[0].weight.data.clone()
visualize_tensor(filter, ch=0, allkernels=False)
```

```
plt.axis('off')
plt.ioff()
plt.title('decoder_conv3 Weights')
plt.show()
```



decoder_conv2 Weights



decoder_conv3 Weights

